



yugabyteDB

Workshop: Deploy a Real-time Polling App with Hasura Cloud and Yugabyte Cloud

Eric Pratt, YugabyteDB


Nikhil Chandrappa, YugabyteDB

Agenda

- Overview of GraphQL Applications using Hasura and YugabyteDB
- Hands on session for Implementing a real-time poll application
 - a. Setup the Hasura + YugabyteDB cloud connectivity
 - b. Running migrations to setup schema for real-time poll application
 - c. Updating node.js application to use Hasura Cloud Instance
 - d. Running the real-poll application
- Tuning and performance of GraphQL Queries
- Do's and Don't for GraphQL queries with YugabyteDB

Getting started with GraphQL

- Provides GraphQL abstractions over your Database
 - a. Query, mutate data via GraphQL constructs
 - b. Build your schema and evolve your domain models
- Evolve your API without versioning
- Provides out of the box pagination and filtering
- Collate disparate Data sources
- Eventing Support




GraphQL Queries for Retrieving Data

Http request

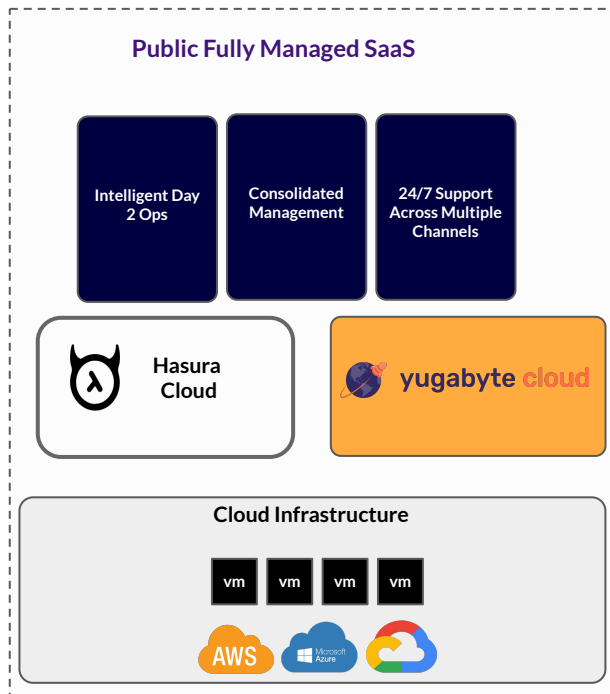
POST /graphql

```
query {  
  author {  
    name  
    articles {  
      title  
      content  
    }  
  }  
}
```



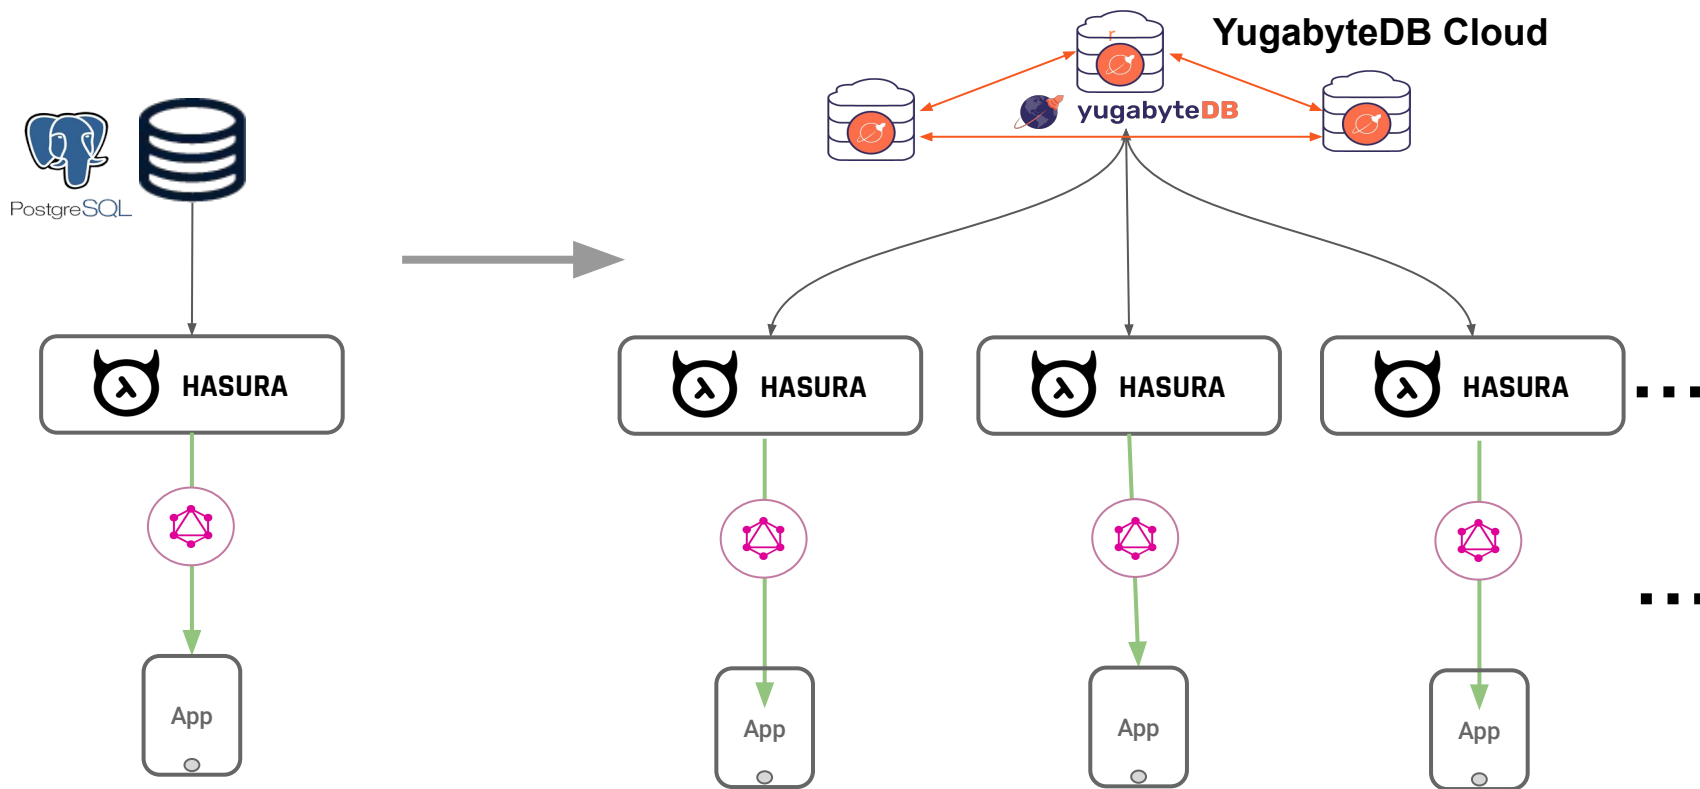
```
{  
  "data": {  
    "author": [  
      {  
        "name": "Prof. John",  
        "articles": [  
          {  
            "title": "John's ..",  
            "content": "incididunt  
ut labore"  
          }  
        ]  
      }  
    ]  
  }  
}
```

Developers embracing cloud



1. Fully managed GraphQL Stack
2. Automated infrastructure operations
3. Removes the need to “install” and “manage” GraphQL engines and backend databases
4. Linear scalability of the backend database at the click of a button

Scaling out GraphQL services with Distributed SQL



Hands-on session

Configuring Hasura Cloud with YugabyteDB Cloud Instance

Hands-on session

Building Real Time poll App using Hasura Cloud and Yugabyte Cloud

Query Tuning - Where to Start?

Performance tuning vs Query debugging

Performance tuning:

- More often involves OS level debugging
- Looking at statistics for memory, cpu, io
- How do queries effect the statistics above?
- Do we have enough horsepower on the DB side to power our applications?
- Identify hotspots or hot nodes


Query debugging and tuning:

- How to identify slow queries in the Database itself
- How to analyze slow queries on Yugabyte platform
- How to analyze slow queries in Hasura
- What to do once we've found a slow query
- Explain vs explain analyze and what it means

How To Analyze Queries With Hasura

Analyzing queries in Hasura

Hasura has a built-in function that allows you to run the equivalent to an explain analyze within Postgres. You can see it from the Hasura console.



Analyze output with Hasura

Query Analysis

Top level nodes

Generated SQL

```
SELECT
  "_subs"."result_id",
  "_fld_resp"."root" AS "result"
FROM
  UNNEST(($ 1) :: uuid [], ($ 2) :: json []) AS "_subs"("result_id", "result_vars")
  LEFT OUTER JOIN LATERAL (
    SELECT
      json_build_object('online_users', "online_users"."root") AS "root"
    FROM
      (
        SELECT
          coalesce(json_agg("root"), '[]') AS "root"
        FROM
          (
            SELECT
              row_to_json(
                (
                  SELECT
                    "_1_e"
```

Execution Plan

```
Nested Loop Left Join  (cost=110.04..112.32 rows=100 width=48)
-> Function Scan on _subs  (cost=0.01..1.00 rows=100 width=16)
-> Materialize  (cost=110.04..110.07 rows=1 width=32)
    -> Subquery Scan on online_users  (cost=110.04..110.06 rows=1 width=32)
        -> Aggregate  (cost=110.04..110.05 rows=1 width=32)
            -> Aggregate  (cost=110.00..110.01 rows=1 width=8)
                -> Seq Scan on "user"  (cost=0.00..107.50 rows=1000 width=0)
                    Filter: (last_seen_at > (now() - '00:00:15'::interval))
            SubPlan 1
                -> Result  (cost=0.00..0.01 rows=1 width=32)
```

How to Identify Slow Queries in YugabyteDB

Finding slow queries in YB

As YugabyteDB is Postgres compatible we can use features built-in to Postgres to dig into problem queries. One of the best options is `pg_stat_statements`.

```
yugabyte=# \d pg_stat_statements
```

View "pg_catalog.pg_stat_statements"


Column	Type	Collation	Nullable	Default
userid	oid			
dbid	oid			
queryid	bigint			
query	text			
calls	bigint			
total_time	double precision			
min_time	double precision			
max_time	double precision			
mean_time	double precision			
stddev_time	double precision			
rows	bigint			


Queries to run against pg_stat_statement


I find the following queries to be the most helpful:


- `SELECT pg_stat_statements_reset();`
 - This allows us to clear out pg_stat_statements
- `select dbid,query,calls,total_time,min_time,max_time,mean_time,rows from pg_stat_statements order by total_time desc limit 10;`


Identifying a slow query in YugabyteDB Platform



Dashboard



Universes


Metrics


Tasks


Alerts


Configs


Help

andrew-test-metrics-ysql-5 Ready

OverviewTablesNodesMetricsQueriesReplicationTasksBackupsHealth

ConnectActions

Live QueriesSlow Queries

Slow Queries

Column Display

- ☒ Query
- ☒ Database
- ☒ User
- ☒ Count
- ☒ Total Time
- ☒ Rows
- ☐ Avg Exec Time (ms)
- ☐ Min Exec Time (ms)
- ☐ Max Exec Time (ms)
- ☐ Std Dev Time
- ☐ Temp Tables RAM

Filter by query text

QUERY

DATABASE

USER

TOTAL TIME

COUNT

ROWS

☐ SELECT a.rolname, t.datname, t.queryid, t.query, t.calls, t.total_time, ... postgres yugabyte 1202.823 178 1779

☐ SELECT n.nspname as "Schema", c.relname as "Name", CASE c.relki... testmetrics yugabyte 121.803 1 1

☐ SELECT pg_catalog.quote_ident(datname) FROM pg_catalog.pg_dat... yugabyte yugabyte 0.450 1 1

☐ SELECT pg_catalog.quote_ident(c.relname) FROM pg_catalog.pg_cl... testmetrics yugabyte 4.129 1 1

☐ BEGIN testmetrics yugabyte 0.004 1 0


☐ SET extra_float_digits = 3 postgres yugabyte 4.365 179 0

☐ ROLLBACK testmetrics yugabyte 0.011 1 0

☐ CREATE TABLE metrics1 (id int primary key, name text) testmetrics yugabyte 295.138 1 0

☐ INSERT INTO metrics1 VALUES (\$1, \$2) testmetrics yugabyte 8.869 2 2

☐ CREATE DATABASE testmetrics yugabyte yugabyte 4446.162 1 0

 yugabyteDB

© 2021 - All Rights Reserved

17

You've Identified a Slow Query, Now What?

Running an explain plan

Let's have a look at an explain plan:

```
EXPLAIN SELECT * FROM foo;
```

QUERY PLAN

```
-----  
Seq Scan on foo (cost=0.00..155.00 rows=10000 width=4)  
(1 row)
```

Here is a basic example taken from the postgres docs on a sample table with a single integer column and 1000 rows.

We can see that this query will do a sequential scan on table foo for all rows.

Understanding explain plan cont.

Let's take a look at the following explain analyze:

```
$ explain analyze select * from t limit 100;  
                QUERY PLAN
```

```
-----  
Limit (cost=0.00..9.33 rows=100 width=608) (actual time=0.008..0.152 rows=100 loops=1)  
-> Seq Scan on t (cost=0.00..93333.86 rows=999986 width=608) (actual time=0.007..0.133  
rows=100 loops=1)  
Total runtime: 0.181 ms  
(3 rows)
```

Similar to explain, explain analyze will show you where exactly the time is spent. In this output we can see that the actual time is recorded in the output.

What about something more complex?

Let's consider the following tables:

```
create table breweries (  
  brew_id int,  
  name varchar(100),  
  city varchar(100),  
  state varchar(3),  
  primary key (brew_id)  
);
```

```
create table beers (  
  name varchar(100),  
  beer_id int,  
  abv real,  
  ibu int,  
  brewery_id int,  
  style varchar(100),  
  ounces real,  
  primary key (beer_id),  
  constraint fk_beer_breweries  
    foreign key (brewery_id)  
    references breweries(brew_id)  
);
```

Let's query those tables and see what it looks like

```
explain
select be.name,
       br.name
from   beers be, breweries br
where  be.brewery_id = br.brew_id;
```

QUERY PLAN

```
-----
Nested Loop  (cost=0.00..213.89 rows=1000 width=436)
-> Seq Scan on beers be  (cost=0.00..100.00 rows=1000 width=222)
-> Index Scan using breweries_pkey on breweries br  (cost=0.00..0.11 rows=1 width=222)
    Index Cond: (brew_id = be.brewery_id)
```



yugabyte**DB**

Thank You

Join us on Slack: yugabyte.com/slack

Star us on Github:

<https://github.com/yugabyte/spring-data-yugabytedb>